

P5 Design Document
Taylor Williamson
125007948
4/16/2020

In this assignment, the task is to implement a scheduler for multiple kernel level threads. A scheduler is a third party whose purpose is to ensure that the threads are playing nice and each one is getting its fair share of CPU time which is determined by the scheduling algorithm. In this assignment we will be implementing a First In First Out (FIFO) scheduler.

The scheduler keeps and maintains a ready queue, which lists threads that are waiting to get to the CPU. When the current thread calls `yield()`, the scheduler gets the head of the ready list and calls the dispatcher to invoke a context switch. The `resume(t1)` method is called when the system decides that thread `t1` should become ready to execute on the CPU again. The `add(t1)` function holds the instructions for making the thread `t1` runnable by the scheduler. This is run after thread creation. Finally, the `terminate(t1)` method removes thread `t1` from the scheduler in preparation for destruction and handles the case where the thread wants to terminate itself.

The threads that kernel.C is producing are non-terminating and only stop executing when they pass control to another thread. When a thread is terminated, we need to worry about releasing the CPU, releasing memory, giving control to the next thread, etc. We will test this by making our kernel threads return after `DEFAULT_NB_ITERS` loops.